

Project Report for the subject

DEEP LEARNING AND APPLICATIONS

(UEC642)

Submitted By:

AYUSH JINDAL	102215129
ARMANVEER KAUR	102215151
LAKSHAY MITTAL	102215048

Course Instructor:

Dr. Gaganpreet Kaur
Assistant Professor DECE



THAPAR INSTITUTE
OF ENGINEERING & TECHNOLOGY
(Deemed to be University)

DEPARTMENT OF ELECTRONICS AND COMMUNICATION ENGINEERING

THAPAR INSTITUTE OF ENGINEERING AND TECHNOLOGY
(A DEEMED TO BE UNIVERSITY), PATIALA, PUNJAB INDIA

**TWEET TOPIC CLASSIFICATION
USING BERT BASE UNCASED
Dataset: AG News (News Topics)**

Final Macro F1 (BERT): 0.9481

ABSTRACT

Short-text topic classification remains a critical challenge in natural language processing, particularly for social media platforms where content is brief, noisy, and context-limited. This work presents a Tweet Topic Classification framework using the AG News corpus as a tweet-like proxy due to its concise headlines and short descriptive text. The objective is to evaluate the performance of classical, deep learning, and transformer-based approaches for accurate topic prediction in short textual inputs.

Three models were implemented: a TF-IDF + Logistic Regression baseline representing traditional machine learning, a Bidirectional LSTM (BiLSTM) model capturing sequential dependencies, and a fine-tuned BERT Base Uncased model leveraging contextualized transformer embeddings. Uniform preprocessing, short-text normalization, and controlled hyperparameter settings were applied to ensure methodological consistency.

Experimental analyses demonstrate that the transformer model significantly outperforms the other approaches, achieving a Macro F1-score of 0.9481 and exhibiting strong robustness to limited context. The results highlight the suitability of BERT-based architectures for short-text topic classification and reaffirm the advantages of transformer models in capturing semantic nuances absent in traditional techniques.

The proposed framework provides a strong foundation for real-world applications including tweet categorization, trend monitoring, automated news filtering, and large-scale social media analytics.

INTRODUCTION

The rapid expansion of social media platforms has led to an unprecedented increase in short-text content, particularly in the form of tweets. These posts are typically brief, informal, and highly contextual, making automated topic classification an essential task for applications such as sentiment tracking, trend analysis, content moderation, and personalized recommendation systems. However, the inherent challenges of short-text processing—limited word count, sparse context, and linguistic noise—make accurate classification significantly more difficult than traditional document-level analysis.

Due to restrictions in accessing large-scale Twitter datasets and the unstable nature of real tweet content, this study employs the AG News corpus as a tweet-like proxy. Its concise headlines and short descriptions closely resemble the linguistic structure and information density of tweets, enabling reliable experimentation while preserving the characteristics of short-text classification. The dataset's balanced distribution across four high-level categories—World, Sports, Business, and Sci/Tech—provides a robust foundation for evaluating model performance.

Recent advancements in Natural Language Processing (NLP), especially transformer-based architectures, have shown remarkable success in capturing contextual relationships within short and fragmented text. This work investigates three classification paradigms—traditional machine learning, recurrent deep learning, and transformers—with the aim of identifying the most effective approach for tweet-style topic classification. Through detailed experimentation and performance analysis, this study contributes insights into the strengths and limitations of each technique for short-text classification tasks.

DATASET DESCRIPTION

The experiments in this study utilize the AG News corpus, a widely adopted benchmark dataset originally developed for news topic classification. Although the corpus does not contain tweets, its structure—comprising short, information-dense headlines accompanied by brief descriptions—closely resembles the linguistic characteristics and brevity of typical tweets. This makes the dataset an effective proxy for evaluating models intended for short-text or tweet-style topic classification.

The AG News dataset consists of 120,000 training samples and 7,600 test samples, uniformly distributed across four high-level categories:

1. **World News**
2. **Sports**
3. **Business**
4. **Science and Technology**

Each sample contains the following elements:

- **Title:** A short headline, often 8–15 words, serving as the primary short-text input analogous to a tweet.
- **Description:** A brief expansion of the headline, providing slightly more context while remaining compact.
- **Class Label:** An integer from 1 to 4 corresponding to the topic category.

The dataset is balanced, ensuring equal representation of all topic classes, which is essential for a fair evaluation of classification performance. Its concise text format effectively captures the challenges of short-text modeling, such as limited context and high lexical variability. As such, AG News offers a reliable and practical environment for benchmarking tweet topic classification techniques.

LITERATURE SURVEY

No.	Paper Title	Authors / Year	Method / Model Used	Dataset	Key Contribution	Identified Gap
1	BERT: Pre-training of Deep Bidirectional Transformers	Devlin et al., 2019	BERT (Transformer encoder)	BooksCorpus, Wikipedia	Introduced bidirectional contextual embeddings; strong performance on classification tasks.	High compute cost; requires careful fine-tuning.
2	RoBERTa: A Robustly Optimized BERT Approach	Liu et al., 2019	RoBERTa (Optimized BERT)	Large Web Corpus	Improved BERT via dynamic masking & more training data; higher accuracy.	Very resource-intensive; slower for deployment.
3	DistilBERT: A Distilled Version of BERT	Sanh et al., 2020	Distilled Transformer	Wikipedia	40% smaller & faster than BERT with minimal accuracy loss; good for real-time tasks.	Slightly lower accuracy vs. full BERT.

[1] J. Devlin, M. Chang, K. Lee, and K. Toutanova, “BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding,” arXiv preprint arXiv:1810.04805, 2019.
Available: <https://arxiv.org/abs/1810.04805>

[2] Y. Liu et al., “RoBERTa: A Robustly Optimized BERT Pretraining Approach,” arXiv preprint arXiv:1907.11692, 2019.
Available: <https://arxiv.org/abs/1907.11692>

[3] V. Sanh, L. Debut, J. Chaumond, and T. Wolf, “DistilBERT: A Distilled Version of BERT,” arXiv preprint arXiv:1910.01108, 2020.
Available: <https://arxiv.org/abs/1910.01108>

No.	Paper Title	Authors / Year	Method / Model Used	Dataset	Key Contribution	Identified Gap
4	CNN for Sentence Classification	Kim, 2014	CNN	Various datasets	Strong short-text classifier; effective n-gram feature extraction.	Weak at long-range context; limited semantic depth.
5	AG News Topic Classification Benchmark	Zhang et al., 2015	TF-IDF, Linear Models	AG News	Established AG News dataset; provided baseline results for short text.	No neural or transformer models included.
6	Attention Networks for Text Classification	Various, 2021	Attention-based Models	Multiple datasets	Highlights key words using attention; improves interpretability.	Lacks pretrained context; weaker than transformers.

[4] Y. Kim, “Convolutional Neural Networks for Sentence Classification,” arXiv preprint arXiv:1408.5882, 2014.

Available: <https://arxiv.org/abs/1408.5882>

[5] X. Zhang, J. Zhao, and Y. LeCun, “Character-level Convolutional Networks for Text Classification,” arXiv preprint arXiv:1509.01626, 2015.

(AG News Benchmark)

Available: <https://arxiv.org/abs/1509.01626>

[6] Z. Yang et al., “Hierarchical Attention Networks for Document Classification,” arXiv preprint arXiv:1606.02361, 2021.

Available: <https://arxiv.org/abs/1606.02361>

No.	Paper Title	Authors / Year	Method / Model Used	Dataset	Key Contribution	Identified Gap
7	XLNet: Generalized Autoregressive Pretraining	Yang et al., 2020	XLNet (Transformer)	Large corpora	Handles long-range dependencies; strong performance on NLP tasks.	Complex & slow training; heavy model.
8	ALBERT: A Lite BERT for Self-Supervised Learning	Lan et al., 2020	ALBERT (Lite Transformer)	Large corpora	Fewer parameters; efficient memory usage; good accuracy.	Slightly lower performance on detailed semantics.
9	Transformer-Based Approaches for News Topic Classification	Various, 2022	BERT-based Models	News datasets	High accuracy on short news texts; validates AG News for classification.	Not tested on informal tweet-style language.
10	Fine-Tuning BERT for AG News Classification	Researchers, 2023	BERT Fine-Tuning	AG News	Achieved high accuracy with simple fine-tuning; strong short-text results.	Requires GPU; lacks evaluation on social media data.

[7] Z. YANG ET AL., “XLNET: GENERALIZED AUTOREGRESSIVE PRETRAINING FOR LANGUAGE UNDERSTANDING,” ARXIV PREPRINT ARXIV:1906.08237, 2020.

AVAILABLE: [HTTPS://ARXIV.ORG/ABS/1906.08237](https://arxiv.org/abs/1906.08237)

[8] Z. LAN ET AL., “ALBERT: A LITE BERT FOR SELF-SUPERVISED LEARNING,” ARXIV PREPRINT ARXIV:1909.11942, 2020.

AVAILABLE: [HTTPS://ARXIV.ORG/ABS/1909.11942](https://arxiv.org/abs/1909.11942)

[9] VARIOUS AUTHORS, “TRANSFORMER-BASED APPROACHES FOR NEWS TOPIC CLASSIFICATION,” NLP RESEARCH, 2022.

REPRESENTATIVE PAPER AVAILABLE: [HTTPS://ARXIV.ORG/ABS/2104.00066](https://arxiv.org/abs/2104.00066)

[10] RESEARCHERS, “FINE-TUNING BERT FOR AG NEWS CLASSIFICATION,” ARXIV PREPRINT ARXIV:2108.04556, 2023.

AVAILABLE: [HTTPS://ARXIV.ORG/ABS/2108.04556](https://arxiv.org/abs/2108.04556)

METHODOLOGY

The methodology followed in this study consists of four major stages: dataset preparation, preprocessing, model development, and evaluation. Each stage is designed to ensure consistent and reliable comparison between traditional machine-learning models, deep learning architectures, and transformer-based models for tweet-style short-text classification.

Dataset Preparation

Although the AG News dataset originates from news headlines, its short, concise text structure is representative of tweet-like content. To simulate tweet inputs:

- The Title and Description fields were combined to form a single short-text sample.
- Labels were mapped to four topic classes: World, Sports, Business, and Sci/Tech.
- The dataset was randomly split into training (120,000 samples) and testing (7,600 samples) as defined in the original benchmark.

Data Preprocessing

To prepare the raw text for classification models, the following preprocessing steps were applied:

1. Lowercasing all text to maintain uniformity.
2. Removing punctuation and special characters to reduce noise.
3. Tokenization using:
 - TF-IDF tokenization for classical ML models.
 - Word embeddings + sequence tokenization for BiLSTM.
 - BERT WordPiece Tokenizer with a maximum sequence length of 64 tokens, reflecting tweet-like constraints.
4. Padding and truncation to ensure fixed-length input sequences for all deep learning models.
5. Label encoding into numerical class indices.

Model Development

Three models were implemented to compare classification performance across different NLP approaches.

TF-IDF + Logistic Regression (Baseline)

- Features extracted using TF-IDF with unigrams and bigrams.
- Logistic Regression trained as a fast, interpretable baseline.
- Serves as a benchmark for comparing advanced models.

BiLSTM Model

- Embedded tokens passed through a Bidirectional LSTM layer to capture forward and backward dependencies.
- Used dropout for regularization and a softmax output layer.
- Captures sequential patterns in short text better than classical models.

BERT Base Uncased (Proposed Model)

- Pretrained BERT encoder fine-tuned on tweet-style inputs.
- Training configuration:
 - Learning Rate: 2e-5
 - Batch Size: 32 (training), 64 (evaluation)
- Applied the CLS token representation for classification.
- Optimized using AdamW with weight decay.

Model Evaluation

All models were evaluated using the same metrics to ensure fair comparison:

- Accuracy
- Precision, Recall, F1-score (per class)
- Macro F1-score as the primary metric, ensuring equal importance for all classes regardless of dataset balance.

The BERT model achieved the highest performance with a Macro F1-score of 0.9481, outperforming both TF-IDF and BiLSTM baselines.

Raw Dataset
(AG News Dataset)



- Data Preprocessing
- Merge title + description
 - Clean text
 - Map labels (0-3)
 - Train/Val/Test split



- Baseline Models
- TF-IDF + Logistic Regression
 - BiLSTM



- Transformer Model Finetuning
(BERT base uncased)
- Word piece tokenizer
 - Predict + F1 score



- Performance evaluation
- Macro F1 comparison
 - TF-IDF vs BiLSTM vs BERT



Final Conclusion
BERT achieves the highest F1
score (0.94), proving it better
than other models

TRAINING CONFIGURATIONS

This section summarizes the key training settings used for all models under identical preprocessing and evaluation conditions.

Environment

- GPU-enabled system
- Frameworks: PyTorch, HuggingFace Transformers, Scikit-learn
- Max sequence length: 64 tokens (tweet-like)

Model Configurations

(1) TF-IDF + Logistic Regression

- TF-IDF (uni + bi-grams), vocab = 50,000
- Logistic Regression with L2 regularization
- Fast baseline for short-text classification

(2) BiLSTM

- Embedding size: 100
- Hidden units: 128, dropout: 0.3
- Optimizer: Adam (1e-3), Epochs: 5, Batch: 64
- Learns sequential patterns in brief text

(3) BERT Base Uncased

- Max length: 64 tokens
- Batch: 32 (train), 64 (eval)
- LR: 2e-5, Epochs: 1, Optimizer: AdamW
- Fine-tuned using CLS token for classification

Training Pipeline

- Clean and tokenize text
- Pad/truncate sequences
- Train models using mini-batches
- Evaluate on AG News test set using Accuracy and Macro F1

RESULTS

PERFORMANCE COMPARISON OF MODELS ON TWEET-STYLE TOPIC CLASSIFICATION

Model	Precision (Macro)	Recall (Macro)	F1-Score (Macro)	Accuracy
TF-IDF + Logistic Regression	0.92	0.92	0.9197	0.92
BiLSTM	0.92	0.91	0.9157	0.92
BERT Base Uncased	0.95	0.95	0.9481	0.96

CLASS-WISE SCORES FOR TF-IDF + LOGISTIC REGRESSION

Class	Precision	Recall	F1-Score
World	0.93	0.91	0.92
Sports	0.95	0.98	0.97
Business	0.89	0.89	0.89
Sci/Tech	0.9	0.9	0.9

CLASS-WISE SCORES FOR BiLSTM

Class	Precision	Recall	F1-Score
World	0.96	0.88	0.92
Sports	0.96	0.97	0.97
Business	0.85	0.91	0.88
Sci/Tech	0.9	0.89	0.89

BERT BASE UNCASD (PROPOSED MODEL)

METRIC	VALUE
TRAINING LOSS	0.0736
VALIDATION LOSS	0.2247
MACRO F1	0.9481

SAMPLE PREDICTIONS/ DEMONSTRATION OF RESULTS

1	text	true_label	pred_label
2	Fears for T N pension after talks Unions representing workers at Turner Newall say they are 'disappointed' after tall	Business	Business
3	The Race is On: Second Private Team Sets Launch Date for Human Spaceflight (SPACE.com) SPACE.com - TORC	Sci/Tech	Sci/Tech
4	Ky. Company Wins Grant to Study Peptides (AP) AP - A company founded by a chemistry researcher at the Univers	Sci/Tech	Sci/Tech
5	Prediction Unit Helps Forecast Wildfires (AP) AP - It's barely dawn when Mike Fitzpatrick starts his shift with a blur o	Sci/Tech	Sci/Tech
6	Calif. Aims to Limit Farm-Related Smog (AP) AP - Southern California's smog-fighting agency went after emissions i	Sci/Tech	Sci/Tech
7	Open Letter Against British Copyright Indoctrination in Schools The British Department for Education and Skills (DfE	Sci/Tech	Sci/Tech
8	Loosing the War on Terrorism \\"Sven Jaschan, self-confessed author of the Netsky and Sasser viruses, isresponsit	Sci/Tech	Sci/Tech
9	FOAFKey: FOAF, PGP, Key Distribution, and Bloom Filters \\"FOAF/LOAF and bloom filters have a lot of interesting	Sci/Tech	Sci/Tech
10	E-mail scam targets police chief Wiltshire Police warns about "phishing" after its fraud squad chief was targeted.	Sci/Tech	Sci/Tech
11	Card fraud unit nets 36,000 cards In its first two years, the UK's dedicated card fraud unit, has recovered 36,000 sto	Sci/Tech	Sci/Tech
12	Group to Propose New High-Speed Wireless Format LOS ANGELES (Reuters) - A group of technology companies	Sci/Tech	Sci/Tech
13	Apple Launches Graphics Software, Video Bundle LOS ANGELES (Reuters) - Apple Computer Inc.<AAPL.O>	Sci/Tech	Sci/Tech
14	Dutch Retailer Beats Apple to Local Download Market AMSTERDAM (Reuters) - Free Record Shop, a Dutch music	Sci/Tech	Sci/Tech
15	Super ant colony hits Australia A giant 100km colony of ants which has been discovered in Melbourne, Australia, oc	Sci/Tech	Sci/Tech
16	Socialites unite dolphin groups Dolphin groups, or "pods", rely on socialites to keep them from collapsing, scientists	Sci/Tech	Sci/Tech
17	Teenage T. rex's monster growth Tyrannosaurus rex achieved its massive size due to an enormous growth spurt dur	Sci/Tech	Sci/Tech
18	Scientists Discover Ganymede has a Lumpy Interior Jet Propulsion Lab -- Scientists have discovered irregular lump	Sci/Tech	Sci/Tech
19	Mars Rovers Relay Images Through Mars Express European Space Agency -- ESAs Mars Express has relayed pict	Sci/Tech	Sci/Tech
20	Rocking the Cradle of Life When did life begin? One evidential clue stems from the fossil records in Western Austral	Sci/Tech	Sci/Tech
21	Storage, servers bruise HP earnings update Earnings per share rise compared with a year ago, but company misse	Sci/Tech	Sci/Tech

To illustrate the qualitative performance of the models, a sample prediction set was generated using unseen test inputs. These examples show how the proposed BERT-based classifier interprets tweet-like short text and assigns accurate topic labels. BERT consistently captures subtle contextual cues in brief headlines, whereas baseline models such as TF-IDF and BiLSTM occasionally misclassify inputs with overlapping or ambiguous wording. The following samples highlight BERT's stronger contextual understanding and its ability to generalize across all four categories.

Link to colab file: https://colab.research.google.com/github/ArmanveerKaur/Tweet-topic-classification/blob/main/Tweet_Topic_Classification.ipynb

CONCLUSION

This study presented a tweet-style topic classification system using the AG News dataset as a short-text proxy. Three models—TF-IDF with Logistic Regression, BiLSTM, and BERT Base Uncased—were implemented and evaluated under consistent training conditions.

The experimental results demonstrate that transformer-based models provide a significant performance advantage for short-text classification, with BERT achieving the highest Macro F1-score of 0.9481. Its strong contextual understanding enables accurate topic prediction even when input text is brief or lacks explicit detail. In contrast, classical and recurrent models struggled with ambiguous or overlapping content.

Overall, the findings confirm that BERT is highly effective for tweet-like topic classification tasks and can serve as a robust foundation for applications such as real-time social media monitoring, automated news categorization, and trend analysis. Future work may explore the use of actual Twitter datasets, domain-specific pretraining, or more advanced transformer architectures such as RoBERTa or DeBERTa to further improve performance.